Lift Yourself Up: Retrieval-augmented Text Generation with Self-Memory

Xin Cheng[1]   **Di Luo**[2]   **Xiuying Chen**[3]   **Lemao Liu**[4]   **Dongyan Zhao**[1]   **Rui Yan**[2]

[1] Peking University   [2] Remin University of China
[3] KAUST   [4] Tencent AI Lab
chengxin1998@stu.pku.edu.cn

Abstract

With direct access to human-written reference as memory, retrieval-augmented generation has achieved much progress in a wide range of text generation tasks. Since better memory would typically prompt better generation (we define this as *primal problem*). The traditional approach for memory retrieval involves selecting memory that exhibits the highest similarity to the input. However, this method is constrained by the quality of the fixed corpus from which memory is retrieved. In this paper, by exploring the duality of the primal problem: better generation also prompts better memory, we propose a novel framework, `Selfmem`, which addresses this limitation by iteratively employing a retrieval-augmented generator to create an unbounded memory pool and using a memory selector to choose one output as memory for the subsequent generation round. This enables the model to leverage its own output, referred to as self-memory, for improved generation. We evaluate the effectiveness of `Selfmem` on three distinct text generation tasks: neural machine translation, abstractive text summarization, and dialogue generation, under two generation paradigms: fine-tuned small model and few-shot LLM. Our approach achieves state-of-the-art results in four directions in `JRC-Acquis`, `XSum` (50.3 ROUGE-1), and `BigPatent` (62.9 ROUGE-1), demonstrating the potential of self-memory in enhancing retrieval-augmented generation models. Furthermore, we conduct thorough analyses of each component in the `Selfmem` framework to identify bottlenecks and provide insights for future research.

1 Introduction

In recent years, retrieval-augmented text generation has garnered increasing interest across various fields such as neural machine translation [1, 22, 14], dialogue response generation [71, 6, 42], language modeling [21, 31, 90], and others. This novel generation paradigm initially equips the generation model, either a fine-tuned small model or a large language model (LLM) with in-context learning capability [5], with access to an external database (typically the training corpus) via information retrieval techniques [63, 33] and then generates text based on the retrieved memory.

In this paradigm, the guiding principle for memory retrieval is to find the memory that exhibits the highest similarity to the current input [31, 84, 45]. This aligns with the human intuition that a more similar demonstration sample typically offers more hints. As demonstrated in Figure 1, for a tunable retrieval-augmented translation model, the memory similarity alone exhibits a strong correlation with the final translation quality, regardless of other factors that may influence translation quality (e.g., polysemy, morphology, and coreference). We define this as the *primal problem*: **better memory prompts better generation**. Consequently, numerous studies have focused on how to retrieve better

memory, ranging from sparse retrieval to dense retrieval [10, 58], from a fixed retriever to a learnable retriever [38, 8], and from sentence-level memory to more fine-grained token-level memory [31, 30].

However, a fundamental limitation exists in all previous works: the memory is retrieved from a fixed corpus and is constrained by the corpus's quality. Due to the finite retrieval space, bounded memory significantly limits the potential of memory-augmented generation models. In this paper, by exploring the *primal problem*'s duality, **better generation also prompts better memory**, we propose a novel framework called `Selfmem`. This framework iteratively employs a retrieval-augmented generator to create an unbounded memory pool and uses a memory selector to choose one output as memory for the subsequent generation round. By combining the *primal* and *dual problem*, a retrieval-augmented generation model can elevate itself using its own output, referred to as self-memory. The key insight behind `Selfmem` is that the text more closely resembling the data distribution during inference is not the training data [76], but the model's own output.


Figure 1: Relation between memory and hypothesis on `JRC-Acquis` En→De dataset. The hypothesis is generated by a retrieval-augmented translator whose memory is retrieved from the training set. The X-axis represents the similarity between memory and the reference.

`Selfmem` consists of two complementary components: a retrieval-augmented generator and a memory selector. The generator operates under two distinct paradigms: fine-tuning a small model or few-shot prompting a LLM. For the former, we supervise the generator with labeled data and retrieved memory, while for the latter, we employ a fixed black-box LLM exclusively for inference alongside retrieved in-context learning samples. We then use the generator's output to train a memory selector based on a specific performance metric. By simply replacing the retrieved memory with unbounded generated memory, we achieve higher-quality generation output (*primal problem*), which subsequently serves as memory for the next round after being refined by the memory selector (*dual problem*).

To evaluate the efficacy of the `Selfmem`, we carry out comprehensive experiments in three distinct text generation tasks: neural machine translation, abstractive text summarization, and dialogue generation. We witness substantial enhancements over robust baselines, attaining state-of-the-art outcomes in `JRC-Acquis` (four directions), `XSum` (50.3 ROUGE-1), and `BigPatent` (62.9 ROUGE-1). To gain deeper insights into the `Selfmem`, we meticulously investigate each crucial component and pinpoint the existing system bottleneck to guide future research endeavors.

2 Related Work

2.1 Retrieval-augmented Text Generation

Since the world is not a snapshot once the training corpus is collected, we can never expect an ever-large model to capture everything in its parameters, even for LLMs like GPT-4 [57]. Therefore, it is crucial to equip these models with an external memory bank to store additional knowledge or useful demonstration examples for solving various NLP tasks[38, 68, 83].

In the translation domain, retrieval techniques have long been employed by the localization industry to enhance human translators' productivity and consistency even before the advent of machine translation [82]. Early works on machine translation primarily focused on utilizing memory for statistical machine translation (SMT) systems [70, 46]. For neural machine translation (NMT), [22] were the first to use search engines to retrieve memory from the training set and incorporate it with an external memory network. Subsequent research explored various aspects of retrieval-augmented NMT, such as memory encoding methods [80, 81, 25], joint training of retrievers and generators with monolingual data [8], memory granularity [30], and memory diversity [14]. For few-shot LLM generation, strategies for in-context example selection have been proposed to improve translation quality [1]. Furthermore, in-context machine translation has been shown to be effective for on-the-fly adaptation [69]. For dialogue response generation tasks, employing exemplar/template retrieval as an intermediate step has proven advantageous for generating informative responses [78, 79, 6, 7]. In-context learning example retrieval also aids in controllable dialogue [42]. Other applications

include abstractive summarization [76, 11, 59], code generation [24], paraphrase generation [29, 73], language modeling [31, 90], counterfactual data generation [18], and question answering [32, 28, 27].

2.2 Neural Text Reranking

By alleviating the discrepancy between training and inference (i.e., exposure bias) and directly optimizing desired metrics, two-stage reranking methods have facilitated significant progress in various text generation tasks. In machine translation, pioneering works by [66] and [56] introduced and popularized discriminative reranking for SMT. In the context of NMT, research has focused on two primary reranking approaches: generative reranking [51, 26, 77] and discriminative reranking [36, 64, 17]. For syntactic parsing, [15] were the first to employ a two-stage reranking method to select outputs from a base parser, while [12] introduced a maximum entropy reranker. In text summarization, RefSum [48] proposed a second-stage summarization framework to address train-test distribution mismatches. SimCLS [49] used pairwise Learning To Rank (LTR) to select candidates with the highest matching scores. SummaReranker [62] adopted a multi-task mixture-of-experts framework to leverage different metrics capturing various aspects of generated candidates. BRIO [50] reused the base model for a second round of fine-tuning with both cross-entropy loss and a candidate-level ranking loss. JGR [67] employed an alternate training paradigm to train the generator and reranker.

A key limitation of these reranking methods is that they represent a one-way process, wherein the selected candidates become the system's final output. In contrast, our framework innovatively utilizes the chosen candidates as memory for the subsequent generation round of a retrieval-augmented generator, which can produce better candidates with enhanced memory.

3   Methods

In this section, we begin with a motivating experiment on *generation as memory* (§ 3.1). Then, we introduce `Selfmem`, a framework comprising a *retrieval-augmented generator* (§ 3.2) and a *memory selector* (§ 3.3). The complete framework and algorithm are illustrated in Figure 2 and Algorithm 1.

3.1   Generation as Memory

The primary motivation behind our framework stems from the observation that the memory, which is more similar in distribution to the data during inference, is not the training data (38.89 BLEU, as shown in Table 1). Instead, it is the model's own output (58.58 BLEU) within the unbounded generation space. One interesting exploration involves directly utilizing the generated output as memory in relation to the *primal problem*: better memory prompts better generation.

We conduct experiments on the `JRC-Acquis` En→De dataset. The first row in Table 1 represents conventional retrieval-augmented training with retrieved memory and achieves a 58.58 BLEU score. However, directly incorporating beam output of this trained model as memory (Beam) back into the generation model does not yield any improvements (row 2), despite its higher similarity to the reference compared to the retrieved ones. We hypothesize two potential reasons for this: (1) the retrieval-augmented generator may not generalize effectively in this context due to the memory distribution shift (from 38.89 to 58.58), and (2) the beam memory does not offer any information gain in comparison to the retrieved memory, even when it exhibits more overlap with the references.

Table 1: Experiments on a fixed retrieval-augmented translator with different memory.

	Memory	Hypothesis
Retrieval	38.89	58.58
Beam	58.58	58.43
Reference	100	90.43
Random	1.14	49.08

To eliminate the first hypothesis, we investigate the best and worst scenarios by using the reference as memory (Reference) and randomly sampled sentences as memory (Random). Table 1 illustrates that a retrieval-augmented generator trained with retrieved memory has already learned to effectively utilize memory information in both oracle and random scenarios, with fixed parameters.

To evaluate the second conjecture, we first define the token sets of the reference, retrieved memory, and beam memory as $\mathcal{R}, \mathcal{M}$, and $\mathcal{B}$, respectively. The overlap token set, denoted by $\mathcal{O}$, is defined as the tokens that overlap with the references in the beam memory but not in the retrieved memory, which is represented as $\mathcal{R} \cap \mathcal{B} - \mathcal{R} \cap \mathcal{M}$. $\mathcal{O}$ is considered as the additional information provided

Figure 2: Overall framework. There are two components in `Selfmem`, a retrieval-augmented generator (a) and a memory selector (b). For the primal problem, (a) takes source and memory as input to generate candidates for (b). For the dual problem, (b) takes as input source and generated candidates to select memory for (a).

by the beam memory. Inspired by the confidence analysis of NMT model [53], we compute the set confidence score, $\psi(\cdot)$, as follows:

$$\psi(\cdot) = \frac{1}{|\cdot|} \sum_{y^i \in \cdot} p(y_i | x, y_{<i}) \tag{1}$$

where $p(y_i | x, y_{<i})$ is defined by the generation model. $\psi(\cdot)$ measures the confidence with which the generation model generates the tokens in the set. The value of $\psi(R)$ is 0.58, while that of $\mathcal{O}$ is 0.76, indicating that the generator is relatively confident in generating tokens in $\mathcal{O}$, and therefore does not need to resort to external memory [35]. Beam search ranks generated candidates based on $p(y|x)$, where the selected memory falls within the confidence region of the generator and consequently provides no information gain. This observation motivates us to select memory according to metrics other than $p(y|x)$ in the memory selector (§3.3).

3.2 Retrieval-augmented Generator

Given a text pair (x, y), where $x = \{\mathrm{x}_1, ..., \mathrm{x}_{|x|}\}$ is the source, $y = \{\mathrm{y}_1, ..., \mathrm{y}_{|y|}\}$ is the target. They could be (document, summary) in summarization, (context, response) in dialogue generation or (source, target) in machine translation. The retrieval-augmented generation would first use x to retrieve memory m from datastore $\mathbb{D}$. Then the generator $G_\xi(x, m)$, parameterized by ξ, would take both x and m as input to generate the target sentence y. In this paper, following standard practice, we choose the training set as $\mathbb{D} = \{(x^i, y^i)\}_{i=1}^{|\mathbb{D}|}$. For LLM as G_ξ, we use the standard in-context learning format to give (x, y) as demonstration example. For tunable generator G_ξ, we only keep the target side of *top-1* retrieval results as memory and we consider two commonly used architectures: **Joint-Encoder** [23, 76, 38] and **Dual-Encoder** [80, 8, 14].

Joint-Encoder   This architecture is the standard encoder-decoder-based model [2, 74]. The input is the concatenation of x and m. The encoder would first map the input into the hidden states H:

$$H = \mathrm{Encoder}(x \, [\mathrm{SEP}] \, m) \tag{2}$$

And the decoder would incorporate H by attention mechanism and generate tokens in an auto-regressive manner:

$$h^i = \mathrm{Decoder}(\mathrm{CrossAttn}(H), y_{<i}) \quad P_{G_\xi}(\cdot | x, y_{<i}) = \mathrm{Softmax}(h^i) \tag{3}$$

Dual-Encoder   Instead of treating x and m as a long sequence, this architecture has two encoders, one for x and the other for m. Their outputs are sequentially attended by the decoder with dual cross

attention as in [14]:

$$H_x = \text{SourceEncoder}(x) \quad H_m = \text{MemoryEncoder}(x) \tag{4}$$

$$h^i = \text{Decoder}(\text{CrossAttn}(H_x, H_m), y_{<i}) \tag{5}$$

We use Transformer [74] as the building block for both architectures and optimize G_ξ with NLL loss:

$$\mathcal{L}_{\text{nll}} = -\sum_{t=1}^{|y|} \log P_{G_\xi}(y_t|x, m, y_{<t}) \tag{6}$$

3.3 Memory Selector

The role of memory selector $S_\theta(x, c)$, parameterized by θ, is to select one candidate c from the candidate pool $\mathbb{C}$ generated by G_ξ based on a specific metric $\Delta(\cdot, \cdot)$. The chosen candidate c is then utilized as memory m for the subsequent generation round of G_ξ. As discussed in §3.1, using $p_{G_\xi}(y|x)$ as the metric $\Delta(\cdot, \cdot)$ would result in falling into the confidence region of G_ξ, leading to no information gain. Moreover, a larger value of $p_{G_\xi}(y|x)$ does not necessarily guarantee improved generation quality [54]. Consequently, we define $\Delta(\cdot, \cdot)$ as model-free metrics that are widely employed for assessing generation quality, such as BLEU for Neural Machine Translation (NMT) and ROUGE for Summarization. Our memory selector takes the concatenation of the source x and candidate c_i as input, and produces a multinomial distribution $p_{S_\theta}(\cdot|x)$ over $\mathbb{C}$.

In this paper, we focus on the role of the memory selector, $S_\theta(x, c)$, which is parameterized by θ. The objective of this selector is to choose a single candidate c from the candidate pool $\mathbb{C}$, generated by G_ξ, based on a specific metric, $\Delta(\cdot, \cdot)$.

$$p_{S_\theta}(c_i|x) = \frac{\exp(S_\theta(x \,[\text{SEP}]\, c_i))}{\sum_{j=1}^{|\mathbb{C}|} \exp(S_\theta(x \,[\text{SEP}]\, c_j))} \tag{7}$$

In accordance with [36], the training goal for S_θ is to minimize the discrepancy between the S_θ's predictions and the scores determined by $\Delta(\cdot, \cdot)$. This divergence is quantified using the Kullback-Leibler (KL) divergence.

$$\mathcal{L}_{\text{kl}} = -\sum_{i=1}^{|\mathbb{C}|} p_M(c_i) \log p_{S_\theta}(c_i|x) \quad \text{where} \quad p_M(c_i) = \frac{\exp(\Delta(c_i, y)/\tau)}{\sum_{j=1}^{|\mathbb{C}|} \exp(\Delta(c_j, y)/\tau)} \tag{8}$$

τ is the temperature to control the smoothness of the distribution. At inference, the output of the S_θ is $\arg\max_{c_i \in \mathbb{C}} p_{S_\theta}(c_i|x)$.

3.4 Combine Generator and Selector

We define two generation modes for G_ξ. The first mode, referred to as the *hypothesis mode*, generates a single output for each input, which is utilized for system evaluation. The second mode, known as the *candidate mode*, produces N outputs for a given input, and is employed for training S_θ as well as memory selection. By integrating two modes together, we present the complete framework of our proposed model, `Selfmem`, as illustrated in Algorithm 1.

4 Experimental Setup

4.1 Dataset

We assess the performance of `Selfmem` on three generation tasks, utilizing a total of seven datasets. **Translation.** We evaluate our framework on `JRC-Acquis` datasets [72], a collection of parallel legislative text of European Union Law. It is the benchmark dataset used in translation memory-augmented NMT task [22, 80, 8, 14]. We choose 4 translation directions, namely, Spanish↔English (Es↔En), German↔English (De↔En). **Summarization.** We evaluate on 2 summarization datasets:

Algorithm 1 `Selfmem` Framework

Require: a dataset $\mathbb{D}$, a retriever R, a memory selection metric $\Delta(\cdot, \cdot)$, a retrieval-augmented generator G_ξ, and a memory selector S_θ
1: retrieve memory $\mathbb{M}$ in $\mathbb{D}$ with R
2: train G_ξ with $\mathbb{D}$ and $\mathbb{M}$ (if not LLM)
3: use G_ξ to generate candidate pool $\mathbb{C}$ with $\mathbb{M}$ in candidate mode
4: train S_θ on $\mathbb{C}$ with $\Delta(\cdot, \cdot)$
5: **while** not converged in the validation set **do**
6:     S_θ selects memory from $\mathbb{C}$ as $\mathbb{M}$
7:     G_ξ generates candidate pool $\mathbb{C}$ with $\mathbb{M}$ in candidate mode
8: **end while**
9: G_ξ generates the final hypothesis with $\mathbb{M}$ in hypothesis mode

---

1) `XSum` [55], extreme summarization, a single-document summarization dataset with highly abstractive articles from British Broadcasting Corporation. 2) `BigPatent` [65], consisting of 1.3 million records of U.S. patent documents along with human-written abstractive summaries. **Dialogue.** We experiment on `DailyDialog` [40], which contains multi-turn dialogs on daily life topics and is used by [13, 3, 88]. The detailed statistics for these datasets can be found in the Appendix.

4.2 Implementation Details

We utilize the BM25 algorithm [63] for retrieval purposes. For all tasks, the candidate generation method consists of beam search with a beam width of 50. The number of iterations is determined by the performance on the validation set. **For translation**, we follow the approach of [81, 8, 14], employing a randomly initialized Transformer$_{\text{base}}$ architecture as G_ξ for trainable small model and XGLM [44] for LLM in-context learning. Evaluation metrics include BLEU, TER, and chrF++ obtained from SACREBLEU[60]. The memory selector S_θ utilizes an XLM-R$_{base}$[16] as backbone, with BLEU serving as $\Delta(\cdot, \cdot)$. **For summarization**, we initialize G_ξ with BART$_{\text{base}}$[37] for `BigPatent` and employ BRIO [50] for `XSum`. The evaluation metric comprises ROUGE (R-1/2/L) [43]. **For dialogue generation**, BART$_{\text{base}}$ serves as the backbone for G_ξ. Our dialogue system is evaluated using BLEU (B-1/2) and Distinct (D-1/2) scores [39]. For both dialogue and summarization tasks, we adhere to the methods of [49, 20], adopting RoBERTa$_{\text{base}}$ [47] as the backbone for S_θ. The linear combination of B-1/2 is chosen as $\Delta(\cdot, \cdot)$ for Dialogue Generation, while R-1/2/L is used for Summarization, following [67]. For further implementation details, please refer to the Appendix. Our code is open sourced[1].

5 Experimental Results

5.1 Machine Translation

We select four translation directions and experiment with two generation paradigms: trainable small models and few-shot prompted LLMs. For trainable models, we explore two architectures (joint and dual, as detailed in §3.2). The baselines comprise two types of translation systems: one being the vanilla sequence-to-sequence model [2, 74] without memory augmentation, and the other consisting of retrieval-augmented translation models focusing on memory encoding [22, 80], memory construction [86], memory retrieval [8], and memory diversity [14]. Based on the experimental results[2] shown in Table 2, `Selfmem` significantly enhances the performance of G_ξ across four translation datasets and two different architectures. This is noteworthy, given that the parameters of the G_ξ remain fixed, with the only variable being the input memory. This finding is consistent with the *primal problem* which posits that improved memory typically leads to better generation results.

The *dual problem* is revealed in Table 3. Self-memory, which essentially represents the model's own output, exhibits greater similarity with the ground truth and serves as a more effective memory for generating the final output. This observation highlights a key distinction between `Selfmem` and

---

[1]`https://github.com/hannibal046/selfmemory`
[2]As higher BLEU scores in this range do not necessarily guarantee a superior translation system [9], we also evaluate our system using TER and chrF++. The results can be found in the Appendix.

Table 2: Results of translation task on `JRC-Acquis` measured by BLEU. Models denoted by the same symbol (⋆ and †) have the same parameters and only differ in memory as input. The bolded numbers show the SOTA performance and the underlined numbers show the second-best result. ∗ denotes the system is significantly better than baselines with *p-value* < 0.05 tested by [34].

System	Es→En		En→Es		De→En		En→De	
	Dev	Test	Dev	Test	Dev	Test	Dev	Test
None Memory								
RNNsearch [2]	55.02	59.34	50.54	50.48	50.20	49.74	44.94	43.98
Transformer [74]	64.08	64.63	62.02	61.80	60.18	60.16	54.65	55.43
Retrieval Memory								
SEG-NMT [22]	60.28	59.34	57.62	57.27	55.63	55.33	49.26	48.80
NMT-pieces [86]	63.97	64.30	61.50	61.56	60.10	60.26	55.54	55.14
G-TFM [80]	66.37	66.21	62.50	62.76	61.85	61.72	57.43	56.88
MonoNMT [8]	67.73	67.42	<u>64.18</u>	63.86	64.48	64.62	58.77	58.42
CMM [14]	67.48	67.76	63.84	64.04	64.22	64.33	58.94	58.69
Transformer$_{dual}$⋆	66.87	67.12	63.14	63.54	64.09	63.36	58.69	58.06
Transformer$_{uni}$†	67.74	67.32	63.93	64.12	64.50	64.40	58.16	58.58
Self-Memory								
Transformer$_{dual}$⋆	**68.63**∗	**69.20**∗	64.12∗	<u>64.67</u>∗	<u>65.06</u>∗	<u>64.98</u>∗	<u>59.26</u>∗	59.49∗
Transformer$_{uni}$†	<u>68.26</u>∗	<u>68.80</u>∗	**66.07**∗	**65.94**∗	**65.32**∗	**65.65**∗	**59.88**∗	**60.11**∗

Table 3: Comparison between retrieval memory and self-memory. The quality of memory and hypothesis is measured by the n-gram overlap with reference (BLEU). All experiments are conducted with Transformer$_{joint}$ on `JRC-Acquis`.

		Retrieval		Self	
		memory	hypothesis	memory	hypothesis
En-De	→	38.89	58.58	57.92	60.11
	←	42.56	64.40	64.32	65.65
En-Es	→	40.67	64.12	63.57	65.94
	←	43.05	67.32	67.78	68.80

previous reranking works [36, 62]. Reranking aims to select candidates of higher quality than the beam output, whereas in `Selfmem`, the chosen candidates serve as memory for the retrieval-augmented generator and do not necessarily need to surpass the quality of the beam hypotheses.

Table 4: Evaluation results of in-context learning with self-memory.

		XGLM-1.7B			XGLM-4.5B			XGLM-7.5B		
		Random	kNN	Self	Random	kNN	Self	Random	kNN	Self
En-De	→	11.51	37.87	40.94	17.51	37.60	38.25	18.48	47.82	48.32
	←	27.42	51.00	51.88	30.62	48.12	48.36	33.03	55.65	55.12
En-Es	→	23.87	46.20	48.56	31.83	48.37	49.17	29.97	53.86	54.32
	←	25.29	51.55	53.13	32.16	48.55	49.22	35.22	57.25	57.56

In Table 4, we present the results of LLM with self-memory. We employ XGLM [44] as our backbone generator, with three different sizes ranging from 1.7B to 7.5B. We utilize the recommended prompt as described in [44]. We select three in-context learning examples and report the average scores from three separate runs, taking into account the sensitivity of example selection in ICL [45]. From the table, we first observe a general trend where few-shot translation performance improves as the size of the model increases. Furthermore, we find that more similar translation demonstrations significantly enhance performance across all model sizes (from random, kNN to Self). This suggests that demonstration examples in in-context learning not only act as triggers for model ability but also

adhere to the *primal problem*, where better demonstration example leads to better generation. Also, by comparing the results in Table 2 and Table 4, we can conclude that the cross-lingual LLM with designed examples still falls short of the supervised baselines in this task.

5.2 Summarization

In this paper, we compare the performance of our trainable model with those of REINA [76], PEGASUS [85], and BART [37]. The results are presented in Table5. Initially, it can be observed that memory has varying impacts on different datasets. The enhancement brought by memory in the `BigPatent` dataset is significantly larger than that in the `XSum` dataset. This can be attributed to the inherent characteristics of the `BigPatent` dataset, which consists of official patent documents that exhibit considerable similarity. Consequently, this greatly improves the summarization quality in accordance with the *primal problem*. Furthermore, we discovered that self-memory substantially enhances the performance of both BRIO (+1.2 R1) and BART (+18.5 R1), achieving state-of-the-art results on both datasets. We selected these baselines for a fair comparison, as they share the same base generator. Due to space constraints, additional comparisons and the confidence region of the SOTA model can be found in the Appendix.

Table 5: Results of summarization task on `XSum` and `BigPatent` measured by ROUGE.

System	Memory	R-1	R-2	R-L		System	Memory	R-1	R-2	R-L
XSum						**BigPatent**				
PEGASUS	None	47.2	24.6	39.3		PEGASUS	None	53.6	33.2	43.2
BRIO	None	49.1	25.6	40.4		BART	None	44.4	21.3	31.0
REINA (PG)	Retrieval	48.2	26.0	40.2		REINA (B)	Retrieval	59.5	42.6	50.6
REINA (B)	Retrieval	43.2	21.0	35.5		REINA (L)	Retrieval	60.7	43.3	51.3
REINA (L)	Retrieval	46.5	24.1	38.6		REINA (PG)	Retrieval	44.6	21.5	33.3
BRIO$_{dual}\star$	Retrieval	48.6	26.1	40.6		BART$_{dual}\star$	Retrieval	57.4	43.3	49.7
BRIO$_{joint}$†	Retrieval	49.5	26.5	41.2		BART$_{joint}$†	Retrieval	59.6	43.4	51.0
BRIO$_{dual}\star$	Self	49.2	26.2	40.8		BART$_{dual}\star$	Self	61.2	44.6	52.3
BRIO$_{joint}$†	Self	**50.3**	**26.7**	**41.6**		BART$_{joint}$†	Self	**62.9**	**48.1**	**59.6**

5.3 Dialogue Generation

As demonstrated in Table 6, the self-memory significantly enhances the performance of the retrieval-augmented generator for dialogue generation tasks. By optimizing memory using BLEU as $\Delta(\cdot, \cdot)$, the self-memory improves the B-1,2 score over retrieved memory by 3.08 B-1 and 0.6 B-2 on BART$_{joint}$. Intriguingly, although `Selfmem` surpasses the baselines in terms of B-1/2, it falls behind in D-1 and D-2, which can be attributed to the trade-off between BLEU score and Distinct score when evaluating a dialogue system [89]. To address this issue, we opt for D-1,2 as $\Delta(\cdot, \cdot)$ when optimizing S_θ, denoted as BART$_{joint}$†(D). The results in Table 6 highlight the remarkable flexibility of `Selfmem` by directly optimizing memory to achieve the desired attributes for diverse and informative dialogue.


(a) Hypothesis

(b) Candidates

Figure 3: (a) shows generation quality in the iteration process with different S_θ in both trainable generator architectures. (b) shows candidates quality in the iteration process with an oracle S_θ.

Table 6: Results of dialogue generation task on `DailyDialog` measured by B-1/2 and D-1/2. $BART_{joint}$ (D) denotes the metric $\Delta(\cdot, \cdot)$ for S_θ is the average of D-1 and D-2.

System	Memory	B-1	B-2	D-1	D-2
NCM [75]	None	33.60	26.80	3.00	12.80
iVAE [19]	None	30.90	24.90	2.90	25.00
PLATO-2 [4]	None	34.80	25.12	3.54	25.11
DialoFlow [41]	None	36.17	27.67	4.56	27.12
BART	None	20.72	11.36	3.92	19.44
$BART_{dual}\star$	Retrieval	29.50	21.89	4.74	26.01
$BART_{joint}\dagger$	Retrieval	36.72	31.55	6.13	35.65
$BART_{dual}\star$	Self	33.43	22.85	4.66	26.16
$BART_{joint}\dagger$	Self	**39.80**	**32.15**	5.84	32.16
$BART_{joint}\dagger$ (D)	Self	36.92	32.09	**9.12**	**37.05**

6   Further Analysis

To gain a deeper insight into `Selfmem`, we first examine the impact of each key component, namely G_ξ and S_θ. Subsequently, we perform a detailed token-level analysis of the generated output concerning their frequency in the training set. Experiments are conducted on the `JRC-Acquis` En→De dataset.

Tuning S_θ   We explored various S_θ by direct selection from the candidate pool based on gold rankings. As shown in Figure 3a, both architectures with enhanced S_θ significantly outperform the current SOTA performance (60.11 BLEU). Moreover, we assessed the candidate pool quality during this iterative process using an oracle S_θ, as displayed in Figure 3b. A clear pattern emerges in this boxplot, revealing improvements in the *oracle*, *quartile*, *average*, and *minimum* scores of the candidate pool. These two experiments jointly clarify the `Selfmem`'s underlying intuition: a retrieval-augmented generator profits from superior memory, which can be chosen from its own unbounded output, and subsequently, the generator with improved memory produces a higher-quality candidate pool for the next selection round. Consequently, the model lift itself up.

Tuning G_ξ   As discussed in §3.1, we demonstrated that a trained retrieval-augmented generator, with fixed parameters, possesses the ability to distinguish between "good" and "bad" memory. This observation not only justifies our decision to maintain a fixed generator within our framework but also implies that the G_ξ is not the current bottleneck of the `Selfmem`.

Frequency Analysis   We conduct a comprehensive token-level analysis by computing the 1-gram F1 scores for generated translations and subsequently categorizing the tokens based on their frequency in the training set. The results are depicted in Figure 4. A noticeable pattern emerges, suggesting that the more frequently a model encounters a token during training,

Figure 4: 1-gram F1 score sorted by training corpus frequency.

the higher the accuracy of the generated output [87]. Moreover, our findings indicate that retrieval-augmented models, particularly those incorporating self-memory augmentation, exhibit superior performance in handling long-tail inputs which are challenges for parametric models [61, 52].

7   Conclusion

For the first time, we investigate the fundamental limitation of bounded memory in the current retrieval-augmented literature. We combine the *primal* and *dual problems* together and propose `Selfmem`, a general framework for retrieval-augmented text generation by uplifting generation model with its own output. We conduct comprehensive experiments across various text generation tasks and different generation paradigms, including trainable small model and few-shot prompted LLM. We surpass strong baselines and improve the state-of-the-art performance in serval datasets. We also

meticulously investigate each crucial component and pinpoint the existing system bottleneck to guide future research endeavors.

References

[1] Sweta Agrawal, Chunting Zhou, Mike Lewis, Luke Zettlemoyer, and Marjan Ghazvininejad. In-context examples selection for machine translation. *CoRR*, abs/2212.02437, 2022.

[2] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. In Yoshua Bengio and Yann LeCun, editors, *3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings*, 2015.

[3] Siqi Bao, Huang He, Fan Wang, Hua Wu, and Haifeng Wang. PLATO: pre-trained dialogue generation model with discrete latent variable. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 85–96. Association for Computational Linguistics, 2020.

[4] Siqi Bao, Huang He, Fan Wang, Hua Wu, Haifeng Wang, Wenquan Wu, Zhen Guo, Zhibin Liu, and Xinchao Xu. PLATO-2: Towards building an open-domain chatbot via curriculum learning. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 2513–2525, Online, August 2021. Association for Computational Linguistics.

[5] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In Hugo Larochelle, Marc'Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin, editors, *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*, 2020.

[6] Deng Cai, Yan Wang, Wei Bi, Zhaopeng Tu, Xiaojiang Liu, Wai Lam, and Shuming Shi. Skeleton-to-response: Dialogue generation guided by retrieval memory. In Jill Burstein, Christy Doran, and Thamar Solorio, editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019, Minneapolis, MN, USA, June 2-7, 2019, Volume 1 (Long and Short Papers)*, pages 1219–1228. Association for Computational Linguistics, 2019.

[7] Deng Cai, Yan Wang, Wei Bi, Zhaopeng Tu, Xiaojiang Liu, and Shuming Shi. Retrieval-guided dialogue response generation via a matching-to-generation framework. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, EMNLP-IJCNLP 2019, Hong Kong, China, November 3-7, 2019*, pages 1866–1875. Association for Computational Linguistics, 2019.

[8] Deng Cai, Yan Wang, Huayang Li, Wai Lam, and Lemao Liu. Neural machine translation with monolingual translation memory. In Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli, editors, *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL/IJCNLP 2021, (Volume 1: Long Papers), Virtual Event, August 1-6, 2021*, pages 7307–7318. Association for Computational Linguistics, 2021.

[9] Chris Callison-Burch, Miles Osborne, and Philipp Koehn. Re-evaluating the role of Bleu in machine translation research. In *11th Conference of the European Chapter of the Association for Computational Linguistics*, pages 249–256, Trento, Italy, April 2006. Association for Computational Linguistics.

[10] Qian Cao and Deyi Xiong. Encoding gated translation memory into neural machine translation. In Ellen Riloff, David Chiang, Julia Hockenmaier, and Jun'ichi Tsujii, editors, *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing, Brussels, Belgium, October 31 - November 4, 2018*, pages 3042–3047. Association for Computational Linguistics, 2018.

[11] Ziqiang Cao, Wenjie Li, Sujian Li, and Furu Wei. Retrieve, rerank and rewrite: Soft template based neural summarization. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 152–161, Melbourne, Australia, July 2018. Association for Computational Linguistics.

[12] Eugene Charniak and Mark Johnson. Coarse-to-fine n-best parsing and maxent discriminative reranking. In Kevin Knight, Hwee Tou Ng, and Kemal Oflazer, editors, *ACL 2005, 43rd Annual Meeting of the Association for Computational Linguistics, Proceedings of the Conference, 25-30 June 2005, University of Michigan, USA*, pages 173–180. The Association for Computer Linguistics, 2005.

[13] Wei Chen, Yeyun Gong, Song Wang, Bolun Yao, Weizhen Qi, Zhongyu Wei, Xiaowu Hu, Bartuer Zhou, Yi Mao, Weizhu Chen, Biao Cheng, and Nan Duan. Dialogved: A pre-trained latent variable encoder-decoder model for dialog response generation. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 4852–4864. Association for Computational Linguistics, 2022.

[14] Xin Cheng, Shen Gao, Lemao Liu, Dongyan Zhao, and Rui Yan. Neural machine translation with contrastive translation memories. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 3591–3601. Association for Computational Linguistics, 2022.

[15] Michael Collins and Terry Koo. Discriminative reranking for natural language parsing. *Comput. Linguistics*, 31(1):25–70, 2005.

[16] Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. Unsupervised cross-lingual representation learning at scale. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 8440–8451. Association for Computational Linguistics, 2020.

[17] Yuntian Deng, Anton Bakhtin, Myle Ott, Arthur Szlam, and Marc'Aurelio Ranzato. Residual energy-based models for text generation. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net, 2020.

[18] Tanay Dixit, Bhargavi Paranjape, Hannaneh Hajishirzi, and Luke Zettlemoyer. CORE: A retrieve-then-edit framework for counterfactual data generation. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Findings of the Association for Computational Linguistics: EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 2964–2984. Association for Computational Linguistics, 2022.

[19] Le Fang, Chunyuan Li, Jianfeng Gao, Wen Dong, and Changyou Chen. Implicit deep latent variable models for text generation. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, EMNLP-IJCNLP 2019, Hong Kong, China, November 3-7, 2019*, pages 3944–3954. Association for Computational Linguistics, 2019.

[20] Jiazhan Feng, Chongyang Tao, Zhen Li, Chang Liu, Tao Shen, and Dongyan Zhao. Reciprocal learning of knowledge retriever and response ranker for knowledge-grounded conversations. In Nicoletta Calzolari, Chu-Ren Huang, Hansaem Kim, James Pustejovsky, Leo Wanner, Key-Sun Choi, Pum-Mo Ryu, Hsin-Hsi Chen, Lucia Donatelli, Heng Ji, Sadao Kurohashi, Patrizia Paggio, Nianwen Xue, Seokhwan Kim, Younggyun Hahm, Zhong He, Tony Kyungil Lee, Enrico Santus,

Francis Bond, and Seung-Hoon Na, editors, *Proceedings of the 29th International Conference on Computational Linguistics, COLING 2022, Gyeongju, Republic of Korea, October 12-17, 2022*, pages 389–399. International Committee on Computational Linguistics, 2022.

[21] Edouard Grave, Armand Joulin, and Nicolas Usunier. Improving neural language modefls with a continuous cache. In *5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24-26, 2017, Conference Track Proceedings*. OpenReview.net, 2017.

[22] Jiatao Gu, Yong Wang, Kyunghyun Cho, and Victor O. K. Li. Search engine guided neural machine translation. In Sheila A. McIlraith and Kilian Q. Weinberger, editors, *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence, (AAAI-18), the 30th innovative Applications of Artificial Intelligence (IAAI-18), and the 8th AAAI Symposium on Educational Advances in Artificial Intelligence (EAAI-18), New Orleans, Louisiana, USA, February 2-7, 2018*, pages 5133–5140. AAAI Press, 2018.

[23] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. Retrieval augmented language model pre-training. In *Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13-18 July 2020, Virtual Event*, volume 119 of *Proceedings of Machine Learning Research*, pages 3929–3938. PMLR, 2020.

[24] Tatsunori B. Hashimoto, Kelvin Guu, Yonatan Oren, and Percy Liang. A retrieve-and-edit framework for predicting structured outputs. In Samy Bengio, Hanna M. Wallach, Hugo Larochelle, Kristen Grauman, Nicolò Cesa-Bianchi, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 31: Annual Conference on Neural Information Processing Systems 2018, NeurIPS 2018, December 3-8, 2018, Montréal, Canada*, pages 10073–10083, 2018.

[25] Qiuxiang He, Guoping Huang, Qu Cui, Li Li, and Lemao Liu. Fast and accurate neural machine translation with translation memory. In Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli, editors, *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL/IJCNLP 2021, (Volume 1: Long Papers), Virtual Event, August 1-6, 2021*, pages 3170–3180. Association for Computational Linguistics, 2021.

[26] Kenji Imamura and Eiichiro Sumita. Ensemble and reranking: Using multiple models in the NICT-2 neural machine translation system at WAT2017. In Toshiaki Nakazawa and Isao Goto, editors, *Proceedings of the 4th Workshop on Asian Translation, WAT@IJCNLP 2017, Taipei, Taiwan, November 27- December 1, 2017*, pages 127–134. Asian Federation of Natural Language Processing, 2017.

[27] Gautier Izacard and Edouard Grave. Leveraging passage retrieval with generative models for open domain question answering. In Paola Merlo, Jörg Tiedemann, and Reut Tsarfaty, editors, *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume, EACL 2021, Online, April 19 - 23, 2021*, pages 874–880. Association for Computational Linguistics, 2021.

[28] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick S. H. Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu, editors, *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, EMNLP 2020, Online, November 16-20, 2020*, pages 6769–6781. Association for Computational Linguistics, 2020.

[29] Amirhossein Kazemnejad, Mohammadreza Salehi, and Mahdieh Soleymani Baghshah. Paraphrase generation by learning how to edit from samples. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 6010–6021. Association for Computational Linguistics, 2020.

[30] Urvashi Khandelwal, Angela Fan, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. Nearest neighbor machine translation. In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*. OpenReview.net, 2021.

[31] Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. Generalization through memorization: Nearest neighbor language models. In *International Conference on Learning Representations*, 2020.

[32] Omar Khattab, Keshav Santhanam, Xiang Lisa Li, David Hall, Percy Liang, Christopher Potts, and Matei Zaharia. Demonstrate-search-predict: Composing retrieval and language models for knowledge-intensive NLP. *CoRR*, abs/2212.14024, 2022.

[33] Omar Khattab and Matei Zaharia. Colbert: Efficient and effective passage search via contextualized late interaction over BERT. In Jimmy X. Huang, Yi Chang, Xueqi Cheng, Jaap Kamps, Vanessa Murdock, Ji-Rong Wen, and Yiqun Liu, editors, *Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval, SIGIR 2020, Virtual Event, China, July 25-30, 2020*, pages 39–48. ACM, 2020.

[34] Philipp Koehn. Statistical significance tests for machine translation evaluation. In *Proceedings of the 2004 Conference on Empirical Methods in Natural Language Processing , EMNLP 2004, A meeting of SIGDAT, a Special Interest Group of the ACL, held in conjunction with ACL 2004, 25-26 July 2004, Barcelona, Spain*, pages 388–395. ACL, 2004.

[35] Ankit Kumar, Ozan Irsoy, Peter Ondruska, Mohit Iyyer, James Bradbury, Ishaan Gulrajani, Victor Zhong, Romain Paulus, and Richard Socher. Ask me anything: Dynamic memory networks for natural language processing. In Maria Florina Balcan and Kilian Q. Weinberger, editors, *Proceedings of The 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pages 1378–1387, New York, New York, USA, 20–22 Jun 2016. PMLR.

[36] Ann Lee, Michael Auli, and Marc'Aurelio Ranzato. Discriminative reranking for neural machine translation. In Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli, editors, *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL/IJCNLP 2021, (Volume 1: Long Papers), Virtual Event, August 1-6, 2021*, pages 7250–7264. Association for Computational Linguistics, 2021.

[37] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. BART: denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 7871–7880. Association for Computational Linguistics, 2020.

[38] Patrick S. H. Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In Hugo Larochelle, Marc'Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin, editors, *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*, 2020.

[39] Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. A diversity-promoting objective function for neural conversation models. In Kevin Knight, Ani Nenkova, and Owen Rambow, editors, *NAACL HLT 2016, The 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, San Diego California, USA, June 12-17, 2016*, pages 110–119. The Association for Computational Linguistics, 2016.

[40] Yanran Li, Hui Su, Xiaoyu Shen, Wenjie Li, Ziqiang Cao, and Shuzi Niu. Dailydialog: A manually labelled multi-turn dialogue dataset. In Greg Kondrak and Taro Watanabe, editors, *Proceedings of the Eighth International Joint Conference on Natural Language Processing, IJCNLP 2017, Taipei, Taiwan, November 27 - December 1, 2017 - Volume 1: Long Papers*, pages 986–995. Asian Federation of Natural Language Processing, 2017.

[41] Zekang Li, Jinchao Zhang, Zhengcong Fei, Yang Feng, and Jie Zhou. Conversations are not flat: Modeling the dynamic information flow across dialogue utterances. In *Proceedings of the 59th*

Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers), pages 128–138, Online, August 2021. Association for Computational Linguistics.

[42] Zekun Li, Wenhu Chen, Shiyang Li, Hong Wang, Jing Qian, and Xifeng Yan. Controllable dialogue simulation with in-context learning. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 4330–4347, Abu Dhabi, United Arab Emirates, December 2022. Association for Computational Linguistics.

[43] Chin-Yew Lin. ROUGE: A package for automatic evaluation of summaries. In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain, July 2004. Association for Computational Linguistics.

[44] Xi Victoria Lin, Todor Mihaylov, Mikel Artetxe, Tianlu Wang, Shuohui Chen, Daniel Simig, Myle Ott, Naman Goyal, Shruti Bhosale, Jingfei Du, Ramakanth Pasunuru, Sam Shleifer, Punit Singh Koura, Vishrav Chaudhary, Brian O'Horo, Jeff Wang, Luke Zettlemoyer, Zornitsa Kozareva, Mona T. Diab, Veselin Stoyanov, and Xian Li. Few-shot learning with multilingual generative language models. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 9019–9052. Association for Computational Linguistics, 2022.

[45] Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. What makes good in-context examples for gpt-3? In Eneko Agirre, Marianna Apidianaki, and Ivan Vulic, editors, *Proceedings of Deep Learning Inside Out: The 3rd Workshop on Knowledge Extraction and Integration for Deep Learning Architectures, DeeLIO@ACL 2022, Dublin, Ireland and Online, May 27, 2022*, pages 100–114. Association for Computational Linguistics, 2022.

[46] Lemao Liu, Hailong Cao, Taro Watanabe, Tiejun Zhao, Mo Yu, and Conghui Zhu. Locally training the log-linear model for SMT. In Jun'ichi Tsujii, James Henderson, and Marius Pasca, editors, *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning, EMNLP-CoNLL 2012, July 12-14, 2012, Jeju Island, Korea*, pages 402–411. ACL, 2012.

[47] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized BERT pretraining approach. *CoRR*, abs/1907.11692, 2019.

[48] Yixin Liu, Zi-Yi Dou, and Pengfei Liu. Refsum: Refactoring neural summarization. In Kristina Toutanova, Anna Rumshisky, Luke Zettlemoyer, Dilek Hakkani-Tür, Iz Beltagy, Steven Bethard, Ryan Cotterell, Tanmoy Chakraborty, and Yichao Zhou, editors, *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2021, Online, June 6-11, 2021*, pages 1437–1448. Association for Computational Linguistics, 2021.

[49] Yixin Liu and Pengfei Liu. Simcls: A simple framework for contrastive learning of abstractive summarization. In Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli, editors, *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL/IJCNLP 2021, (Volume 2: Short Papers), Virtual Event, August 1-6, 2021*, pages 1065–1072. Association for Computational Linguistics, 2021.

[50] Yixin Liu, Pengfei Liu, Dragomir R. Radev, and Graham Neubig. BRIO: bringing order to abstractive summarization. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 2890–2903. Association for Computational Linguistics, 2022.

[51] Yuchen Liu, Long Zhou, Yining Wang, Yang Zhao, Jiajun Zhang, and Chengqing Zong. A comparable study on model averaging, ensembling and reranking in NMT. In Min Zhang, Vincent Ng, Dongyan Zhao, Sujian Li, and Hongying Zan, editors, *Natural Language Processing*

and Chinese Computing - 7th CCF International Conference, NLPCC 2018, Hohhot, China, August 26-30, 2018, Proceedings, Part II, volume 11109 of *Lecture Notes in Computer Science*, pages 299–308. Springer, 2018.

[52] Alexander Long, Wei Yin, Thalaiyasingam Ajanthan, Vu Nguyen, Pulak Purkait, Ravi Garg, Alan Blair, Chunhua Shen, and Anton van den Hengel. Retrieval augmented classification for long-tail visual recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6959–6969, 2022.

[53] Yu Lu, Jiali Zeng, Jiajun Zhang, Shuangzhi Wu, and Mu Li. Learning confidence for transformer-based neural machine translation. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 2353–2364. Association for Computational Linguistics, 2022.

[54] Clara Meister, Ryan Cotterell, and Tim Vieira. If beam search is the answer, what was the question? In Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu, editors, *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, EMNLP 2020, Online, November 16-20, 2020*, pages 2173–2185. Association for Computational Linguistics, 2020.

[55] Shashi Narayan, Shay B. Cohen, and Mirella Lapata. Don't give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1797–1807, Brussels, Belgium, October-November 2018. Association for Computational Linguistics.

[56] Franz Josef Och, Daniel Gildea, Sanjeev Khudanpur, Anoop Sarkar, Kenji Yamada, Alexander M. Fraser, Shankar Kumar, Libin Shen, David Smith, Katherine Eng, Viren Jain, Zhen Jin, and Dragomir R. Radev. A smorgasbord of features for statistical machine translation. In Julia Hirschberg, Susan T. Dumais, Daniel Marcu, and Salim Roukos, editors, *Human Language Technology Conference of the North American Chapter of the Association for Computational Linguistics, HLT-NAACL 2004, Boston, Massachusetts, USA, May 2-7, 2004*, pages 161–168. The Association for Computational Linguistics, 2004.

[57] OpenAI. GPT-4 technical report. *CoRR*, abs/2303.08774, 2023.

[58] Md. Rizwan Parvez, Wasi Uddin Ahmad, Saikat Chakraborty, Baishakhi Ray, and Kai-Wei Chang. Retrieval augmented code generation and summarization. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Findings of the Association for Computational Linguistics: EMNLP 2021, Virtual Event / Punta Cana, Dominican Republic, 16-20 November, 2021*, pages 2719–2734. Association for Computational Linguistics, 2021.

[59] Hao Peng, Ankur P. Parikh, Manaal Faruqui, Bhuwan Dhingra, and Dipanjan Das. Text generation with exemplar-based adaptive decoding. In Jill Burstein, Christy Doran, and Thamar Solorio, editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019, Minneapolis, MN, USA, June 2-7, 2019, Volume 1 (Long and Short Papers)*, pages 2555–2565. Association for Computational Linguistics, 2019.

[60] Matt Post. A call for clarity in reporting BLEU scores. In Ondrej Bojar, Rajen Chatterjee, Christian Federmann, Mark Fishel, Yvette Graham, Barry Haddow, Matthias Huck, Antonio Jimeno-Yepes, Philipp Koehn, Christof Monz, Matteo Negri, Aurélie Névéol, Mariana L. Neves, Matt Post, Lucia Specia, Marco Turchi, and Karin Verspoor, editors, *Proceedings of the Third Conference on Machine Translation: Research Papers, WMT 2018, Belgium, Brussels, October 31 - November 1, 2018*, pages 186–191. Association for Computational Linguistics, 2018.

[61] Vikas Raunak, Siddharth Dalmia, Vivek Gupta, and Florian Metze. On long-tailed phenomena in neural machine translation. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3088–3095, Online, November 2020. Association for Computational Linguistics.

[62] Mathieu Ravaut, Shafiq R. Joty, and Nancy F. Chen. Summareranker: A multi-task mixture-of-experts re-ranking framework for abstractive summarization. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 4504–4524. Association for Computational Linguistics, 2022.

[63] Stephen E. Robertson and Hugo Zaragoza. The probabilistic relevance framework: BM25 and beyond. *Found. Trends Inf. Retr.*, 3(4):333–389, 2009.

[64] Julian Salazar, Davis Liang, Toan Q. Nguyen, and Katrin Kirchhoff. Masked language model scoring. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 2699–2712. Association for Computational Linguistics, 2020.

[65] Eva Sharma, Chen Li, and Lu Wang. BIGPATENT: A large-scale dataset for abstractive and coherent summarization. In Anna Korhonen, David R. Traum, and Lluís Màrquez, editors, *Proceedings of the 57th Conference of the Association for Computational Linguistics, ACL 2019, Florence, Italy, July 28- August 2, 2019, Volume 1: Long Papers*, pages 2204–2213. Association for Computational Linguistics, 2019.

[66] Libin Shen, Anoop Sarkar, and Franz Josef Och. Discriminative reranking for machine translation. In Julia Hirschberg, Susan T. Dumais, Daniel Marcu, and Salim Roukos, editors, *Human Language Technology Conference of the North American Chapter of the Association for Computational Linguistics, HLT-NAACL 2004, Boston, Massachusetts, USA, May 2-7, 2004*, pages 177–184. The Association for Computational Linguistics, 2004.

[67] Weizhou Shen, Yeyun Gong, Yelong Shen, Song Wang, Xiaojun Quan, Nan Duan, and Weizhu Chen. Joint generator-ranker learning for natural language generation. *CoRR*, abs/2206.13974, 2022.

[68] Weijia Shi, Sewon Min, Michihiro Yasunaga, Minjoon Seo, Rich James, Mike Lewis, Luke Zettlemoyer, and Wen-tau Yih. REPLUG: retrieval-augmented black-box language models. *CoRR*, abs/2301.12652, 2023.

[69] Suzanna Sia and Kevin Duh. In-context learning as maintaining coherency: A study of on-the-fly machine translation using large language models. *CoRR*, abs/2305.03573, 2023.

[70] Michel Simard and Pierre Isabelle. Phrase-based machine translation in a computer-assisted translation environment. In *Proceedings of Machine Translation Summit XII: Papers, MTSummit 2009, Ottawa, Canada, August 26-30, 2009*, 2009.

[71] Yiping Song, Rui Yan, Xiang Li, Dongyan Zhao, and Ming Zhang. Two are better than one: An ensemble of retrieval- and generation-based dialog systems. *CoRR*, abs/1610.07149, 2016.

[72] Ralf Steinberger, Bruno Pouliquen, Anna Widiger, Camelia Ignat, Tomaz Erjavec, Dan Tufis, and Dániel Varga. The jrc-acquis: A multilingual aligned parallel corpus with 20+ languages. In Nicoletta Calzolari, Khalid Choukri, Aldo Gangemi, Bente Maegaard, Joseph Mariani, Jan Odijk, and Daniel Tapias, editors, *Proceedings of the Fifth International Conference on Language Resources and Evaluation, LREC 2006, Genoa, Italy, May 22-28, 2006*, pages 2142–2147. European Language Resources Association (ELRA), 2006.

[73] Yixuan Su, David Vandyke, Simon Baker, Yan Wang, and Nigel Collier. Keep the primary, rewrite the secondary: A two-stage approach for paraphrase generation. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 560–569, Online, August 2021. Association for Computational Linguistics.

[74] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pages 5998–6008, 2017.

[75] Oriol Vinyals and Quoc V. Le. A neural conversational model. *CoRR*, abs/1506.05869, 2015.

[76] Shuohang Wang, Yichong Xu, Yuwei Fang, Yang Liu, Siqi Sun, Ruochen Xu, Chenguang Zhu, and Michael Zeng. Training data is more valuable than you think: A simple and effective method by retrieving from training data. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 3170–3179. Association for Computational Linguistics, 2022.

[77] Yuguang Wang, Shanbo Cheng, Liyang Jiang, Jiajun Yang, Wei Chen, Muze Li, Lin Shi, Yanfeng Wang, and Hongtao Yang. Sogou neural machine translation systems for WMT17. In Ondrej Bojar, Christian Buck, Rajen Chatterjee, Christian Federmann, Yvette Graham, Barry Haddow, Matthias Huck, Antonio Jimeno-Yepes, Philipp Koehn, and Julia Kreutzer, editors, *Proceedings of the Second Conference on Machine Translation, WMT 2017, Copenhagen, Denmark, September 7-8, 2017*, pages 410–415. Association for Computational Linguistics, 2017.

[78] Jason Weston, Emily Dinan, and Alexander H. Miller. Retrieve and refine: Improved sequence generation models for dialogue. In Aleksandr Chuklin, Jeff Dalton, Julia Kiseleva, Alexey Borisov, and Mikhail S. Burtsev, editors, *Proceedings of the 2nd International Workshop on Search-Oriented Conversational AI, SCAI@EMNLP 2018, Brussels, Belgium, October 31, 2018*, pages 87–92. Association for Computational Linguistics, 2018.

[79] Yu Wu, Furu Wei, Shaohan Huang, Yunli Wang, Zhoujun Li, and Ming Zhou. Response generation by context-aware prototype editing. In *The Thirty-Third AAAI Conference on Artificial Intelligence, AAAI 2019, The Thirty-First Innovative Applications of Artificial Intelligence Conference, IAAI 2019, The Ninth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2019, Honolulu, Hawaii, USA, January 27 - February 1, 2019*, pages 7281–7288. AAAI Press, 2019.

[80] Mengzhou Xia, Guoping Huang, Lemao Liu, and Shuming Shi. Graph based translation memory for neural machine translation. In *The Thirty-Third AAAI Conference on Artificial Intelligence, AAAI 2019, The Thirty-First Innovative Applications of Artificial Intelligence Conference, IAAI 2019, The Ninth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2019, Honolulu, Hawaii, USA, January 27 - February 1, 2019*, pages 7297–7304. AAAI Press, 2019.

[81] Jitao Xu, Josep Maria Crego, and Jean Senellart. Boosting neural machine translation with similar translations. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 1580–1590. Association for Computational Linguistics, 2020.

[82] Masaru Yamada. The effect of translation memory databases on productivity. *Translation research projects*, 3:63–73, 2011.

[83] Michihiro Yasunaga, Armen Aghajanyan, Weijia Shi, Rich James, Jure Leskovec, Percy Liang, Mike Lewis, Luke Zettlemoyer, and Wen-tau Yih. Retrieval-augmented multimodal language modeling. *CoRR*, abs/2211.12561, 2022.

[84] Dani Yogatama, Cyprien de Masson d'Autume, and Lingpeng Kong. Adaptive semiparametric language models. *Trans. Assoc. Comput. Linguistics*, 9:362–373, 2021.

[85] Jingqing Zhang, Yao Zhao, Mohammad Saleh, and Peter J. Liu. PEGASUS: pre-training with extracted gap-sentences for abstractive summarization. In *Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13-18 July 2020, Virtual Event*, volume 119 of *Proceedings of Machine Learning Research*, pages 11328–11339. PMLR, 2020.

[86] Jingyi Zhang, Masao Utiyama, Eiichiro Sumita, Graham Neubig, and Satoshi Nakamura. Guiding neural machine translation with retrieved translation pieces. In Marilyn A. Walker, Heng Ji, and Amanda Stent, editors, *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2018, New Orleans, Louisiana, USA, June 1-6, 2018, Volume 1 (Long Papers)*, pages 1325–1335. Association for Computational Linguistics, 2018.

[87] Tong Zhang, Wei Ye, Baosong Yang, Long Zhang, Xingzhang Ren, Dayiheng Liu, Jinan Sun, Shikun Zhang, Haibo Zhang, and Wen Zhao. Frequency-aware contrastive learning for neural machine translation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 11712–11720, 2022.

[88] Xueliang Zhao, Lemao Liu, Tingchen Fu, Shuming Shi, Dongyan Zhao, and Rui Yan. Towards efficient dialogue pre-training with transferable and interpretable latent structure. *CoRR*, abs/2210.12461, 2022.

[89] Yinhe Zheng, Zikai Chen, Rongsheng Zhang, Shilei Huang, Xiaoxi Mao, and Minlie Huang. Stylized dialogue response generation using stylized unpaired texts. In *Thirty-Fifth AAAI Conference on Artificial Intelligence, AAAI 2021, Thirty-Third Conference on Innovative Applications of Artificial Intelligence, IAAI 2021, The Eleventh Symposium on Educational Advances in Artificial Intelligence, EAAI 2021, Virtual Event, February 2-9, 2021*, pages 14558–14567. AAAI Press, 2021.

[90] Zexuan Zhong, Tao Lei, and Danqi Chen. Training language models with memory augmentation. *CoRR*, abs/2205.12674, 2022.